



Docker Deep Dive

Daniel Klopp

The Talk

- I'm not telling you what fishing rod to use

The Talk

- I'm not telling you what fishing rod to use
- I'm helping you understand the fishing rod

The Talk

- I'm not telling you what fishing rod to use
- I'm helping you understand the fishing rod
- You can decide if you should fish with it

The Talk

- This is an intermediate deep dive.

The Talk

- This is an intermediate deep dive.
 - Don't panic

The Talk

- This is an intermediate deep dive.
 - Don't panic
- Each category starts easy, and assumes progressively more knowledge

The Talk

- This is an intermediate deep dive.
 - Don't panic
- Each category starts easy, and assumes progressively more knowledge
 - Don't panic!

The Talk

- This is an intermediate deep dive.
 - Don't panic
- Each category starts easy, and assumes progressively more knowledge
 - Don't panic!
- This is not a demo

The Talk

- This is an intermediate deep dive.
 - Don't panic
- Each category starts easy, and assumes progressively more knowledge
 - Don't panic!
- This is not a demo
- If you have questions, ask.

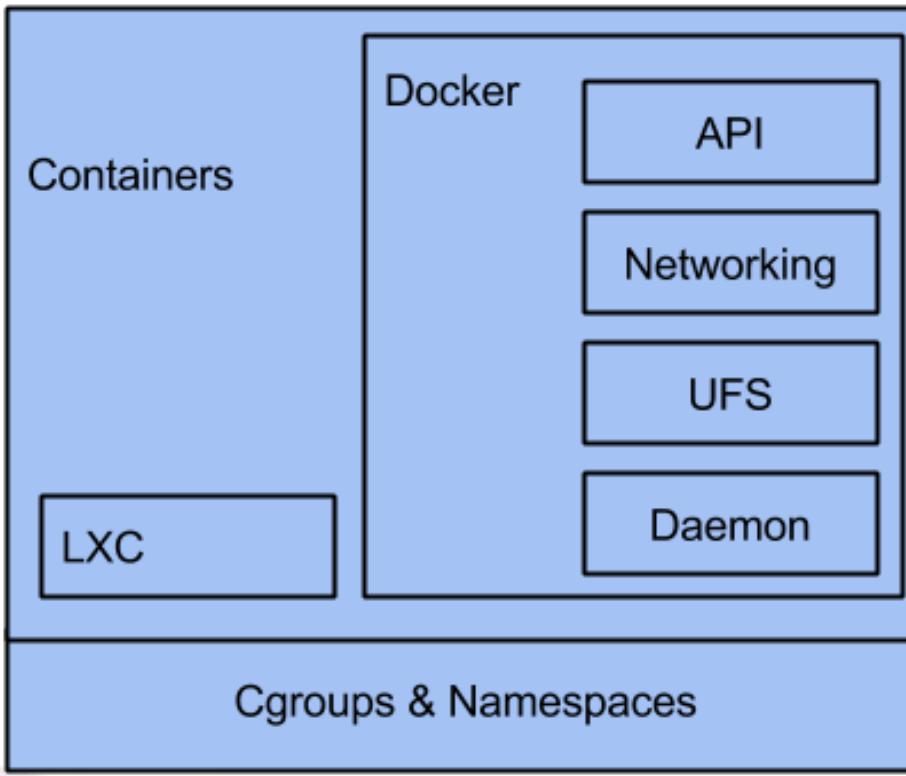
Assumptions

- I assume basic Linux knowledge
- I assume basic TCP/IP knowledge
- I assume basic Filesystem layout knowledge
- No knowledge of containers necessary

Who Am I?

- EE by training, DevOps by practice
- Worked in HPC at National Radio Astronomy Observatory
 - Including kernel development
- Taos, Unix / Scripting / DevOps handiman
- Unix Practice Leader at Taos Consulting

The Talk



Outline

- What are containers
- Containers, hypervisors, jails, chroot
- Docker & Containers Intro
- Linux Container Implementation
 - LXC

Outline

- Docker
- Docker Filesystem
- Docker Networking
- Docker API
- Docker Management Tools

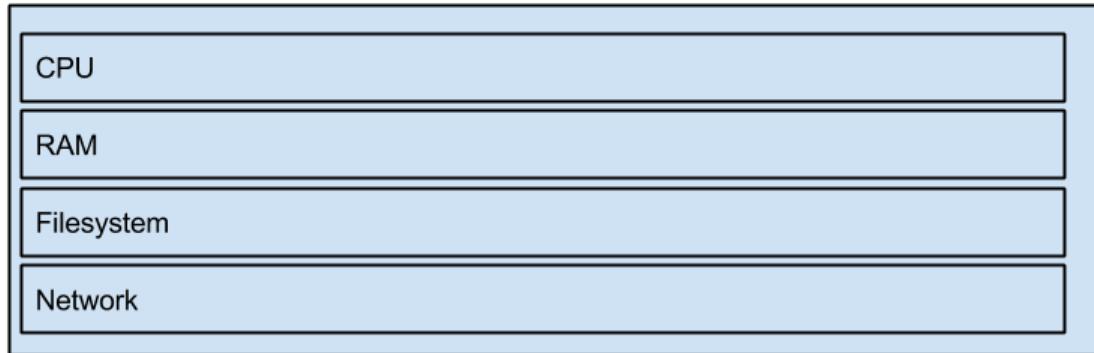
Section

What Are Containers?

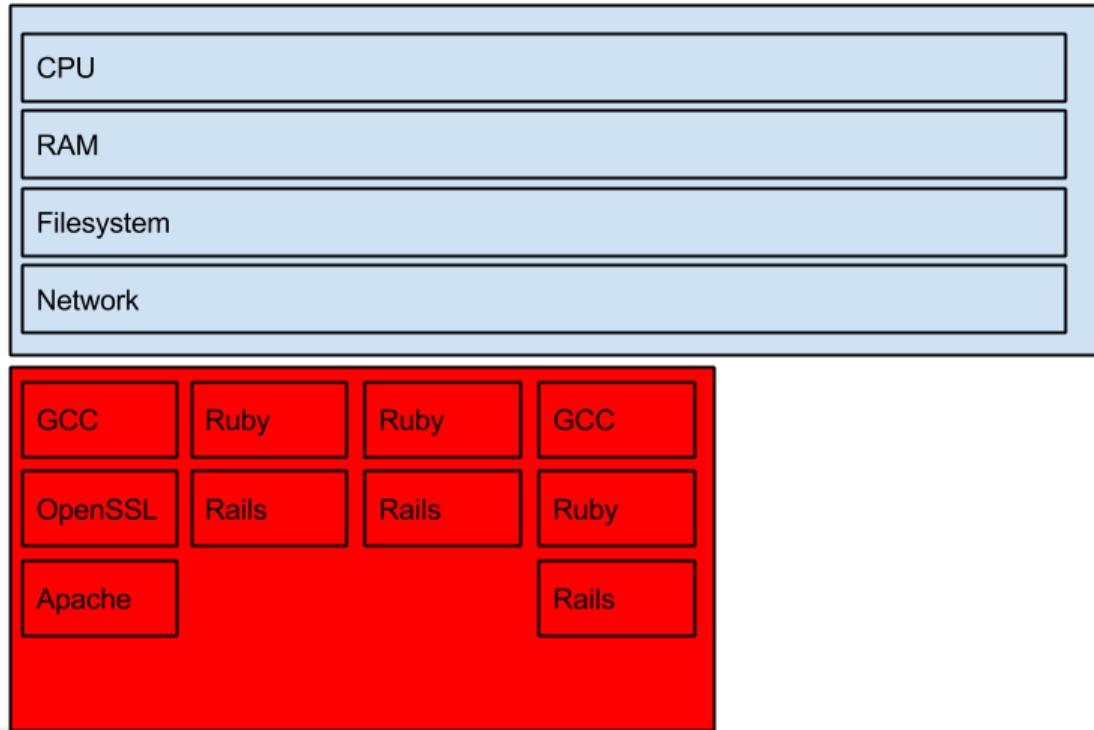
- A brief story
- Containers

A Story

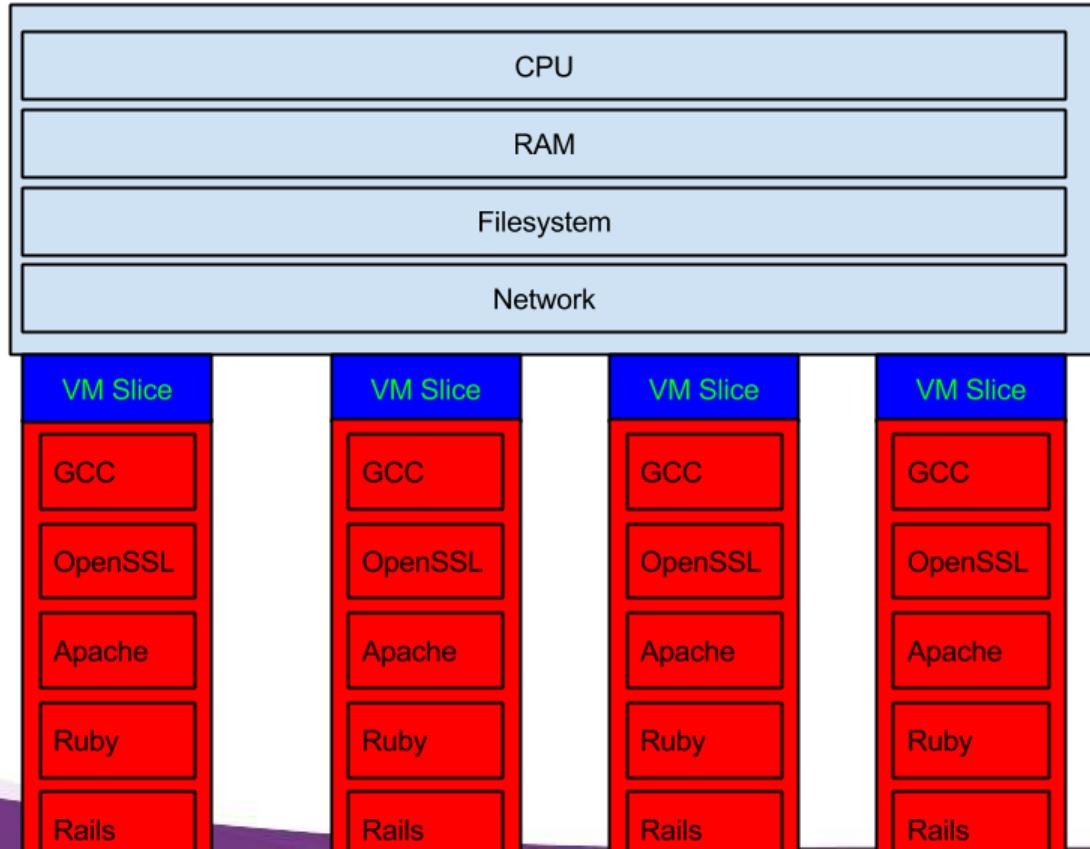
- The good ol' days



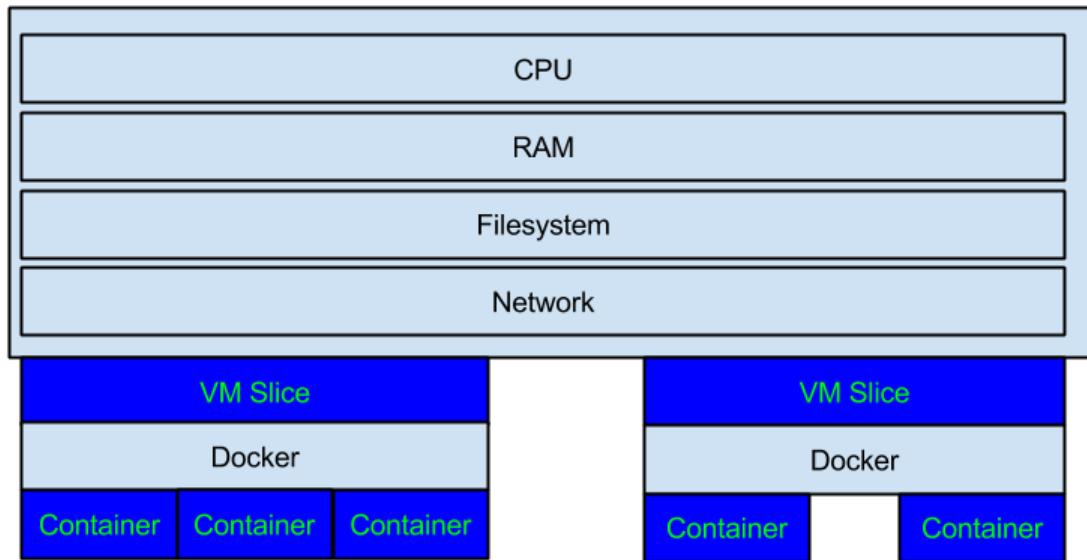
A Story



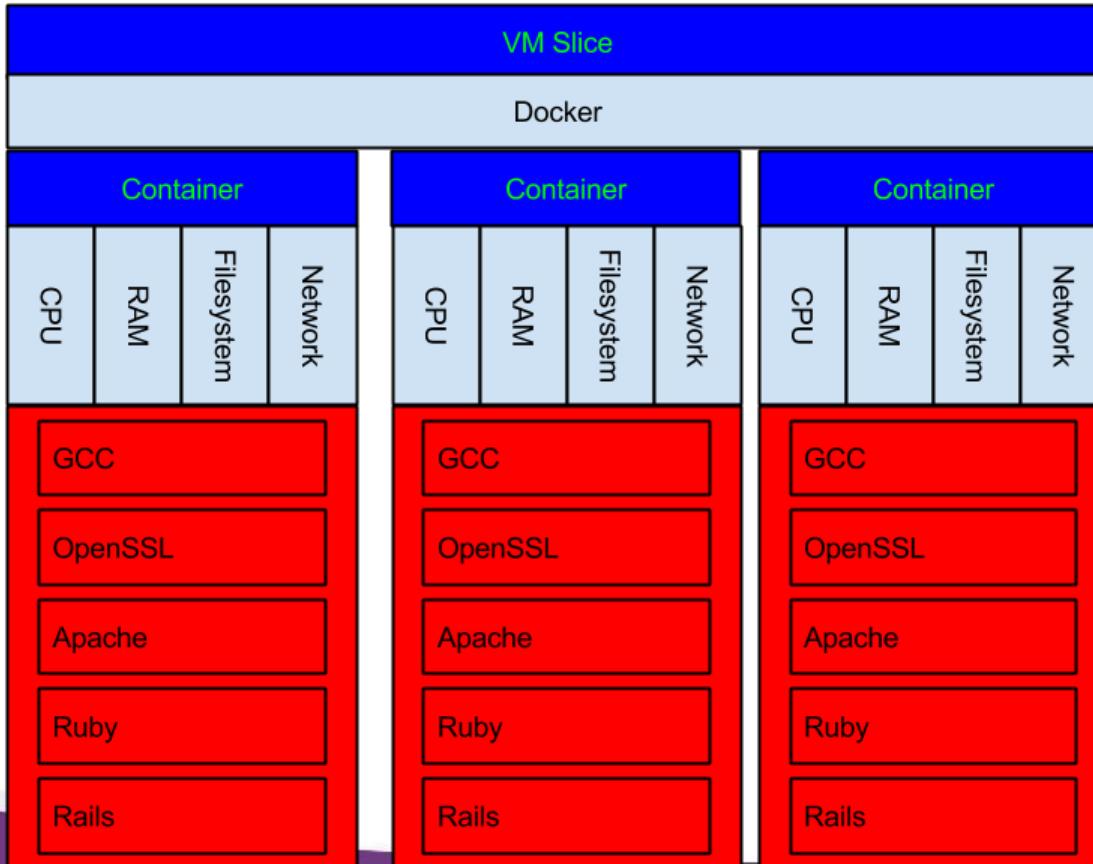
A Story



A Story



A Story



How Is This Achieved?

- Namely
 - Linux Namespaces
 - Kernel cgroups
 - Union Filesystems

How Is This Achieved?

- Namely
 - Linux Namespaces
 - Kernel cgroups
 - Union Filesystems
- We'll review this foundation soon

Kernel Cgroups

- Started by Google
 - Paul Menage and Rohit Seth
- Merged in kernel 2.6.24 (2008)
- Provides
 - Resource Limitation
 - Prioritization
 - Accounting
 - Control

Linux Namespaces

- Isolation
 - Deployment
 - Security

Linux Namespaces

- Isolation
 - Deployment
 - Security
- We will hardly cover this.

Section

Containers, Hypervisors, Jails

- Operating Systems
- Hypervisors
- Jails
- Chroots
- Containers

Operating Systems

- Manage hardware resources
- Provide shared interfaces for programs
- Police programs
- Operating Systems
 - Android
 - Linux
 - Windows
 - BSD
 - Mac

Hypervisors

- Hypervisors virtualize and manage hardware resources
- Popular Type I Hypervisors:
 - KVM
 - Xen
 - VMWare ESX(i)
- Popular Type II Hypervisors:
 - Virtualbox

Chroot

- In practice used to narrow filesystem view
- Cannot isolate process tree
- Cannot enforce resource limits
 - Filesystem only exception
- 20 customers, 20 chroots?
 - 1 misbehaving customer, all customers suffer
- Suffers from escalation insecurities

Jails

- BSD and Solaris
- Operating System level container

Containers

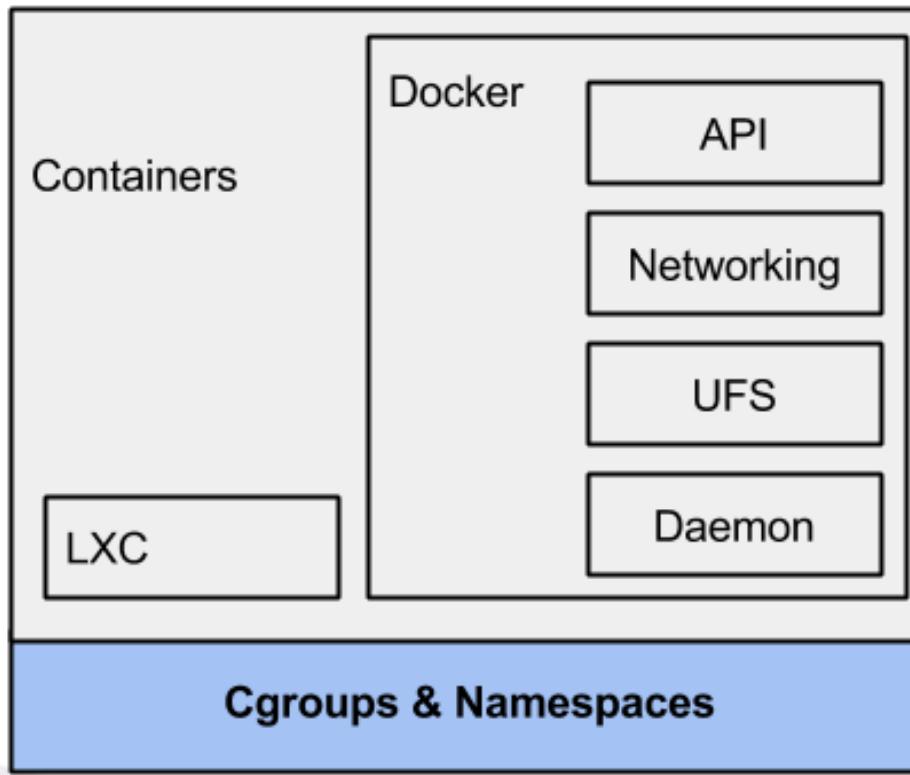
- Operating System level container
- Almost, but not quite, a Linux “jail”
 - (Don’t kill the messenger!)
- Provides meaningful resource controls
 - CPU
 - RAM
- Provides isolation

Section

Linux Container Implementation

- Cgroups
- Namespaces

Where We Are

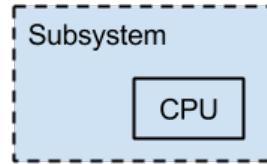


Kernel Cgroups

- Hierarchical

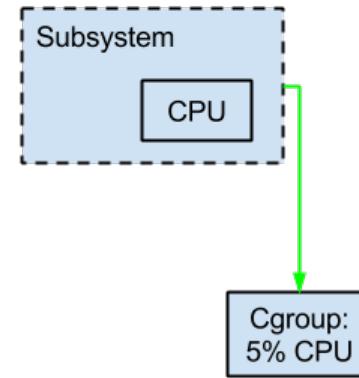
Kernel Cgroups

- Hierarchical
 - Top-level are “subsystems” or “controllers”



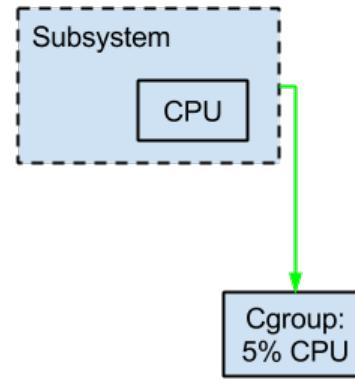
Kernel Cgroups

- Hierarchical
 - Top-level are “subsystems” or “controllers”
 - Lower-level are “cgroups”



Kernel Cgroups

- Hierarchical
 - Top-level are “subsystems” or “controllers”
 - Lower-level are “cgroups”
 - “cgroups” can be nested



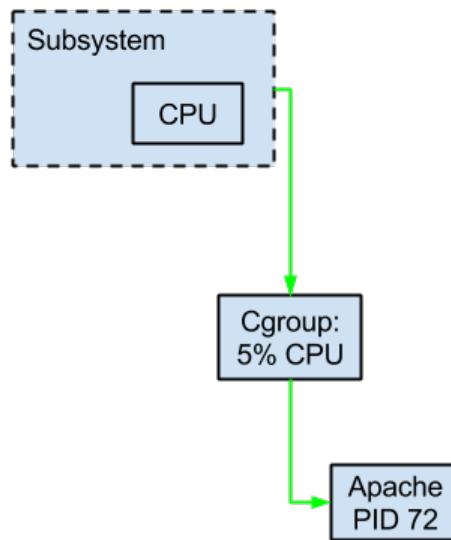
Subsystems

- Also known as “Controllers”
- They are the root of a hierarchy
- They control resources, such as
 - CPU
 - Memory
 - Network
- A “cgroup” is attached to a subsystem, forming a hierarchy

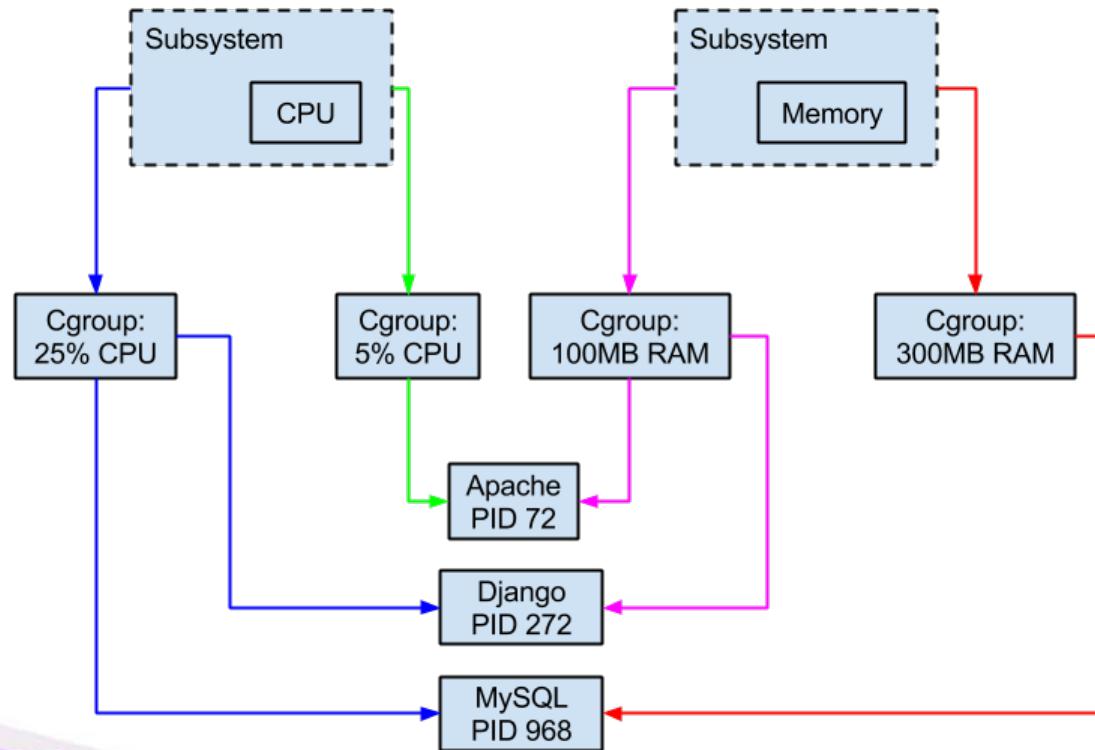
Cgroups

- Encapsulate a set of tasks
- Encapsulate a set of parameters
- The parameters confine the tasks
 - Limiting memory, cpu shares, etc
- Tasks are processes in Linux parlance

Simple Cgroup Hierarchy



Complicated Cgroup Hierarchy



Interacting with Cgroups

- View subsystems in `/sys/fs/cgroup`
 - They are mounted cgroup filesystems

Interacting with Cgroups

- View subsystems in `/sys/fs/cgroup`
 - They are mounted cgroup filesystems
- Memory root cgroup: `/sys/fs/cgroup/memory`
- CPU root cgroup in: `/sys/fs/cgroup/cpu`

Interacting with Cgroups

- View subsystems in `/sys/fs/cgroup`
 - They are mounted cgroup filesystems
- Memory root cgroup: `/sys/fs/cgroup/memory`
- CPU root cgroup in: `/sys/fs/cgroup/cpu`
- They are filesystems, view subsystems with 'mount' command

Some Cgroups on Fedora 21

- Sample mounted cgroups (output formatted)

```
cgroup on /sys/fs/cgroup/cpuset type cgroup (rw,nosuid,nodev,noexec,relatime,cpuset)
```

```
cgroup on /sys/fs/cgroup/cpu,cpuacct type cgroup (rw,nosuid,nodev,noexec,relatime,cpu,cpuacct)
```

```
cgroup on /sys/fs/cgroup/memory type cgroup (rw,nosuid,nodev,noexec,relatime,memory)
```

Interacting with Cgroups

- Let's write a horrible python program

```
def eat_all_ram():
    x=[]
    for i in itertools.count(start=0, step=1):
        x.append(operation(i))
    return len(x)
eat_all_ram()
```

Interacting with Cgroups

- Let's write a horrible python program

```
def eat_all_ram():
    x=[]
    for i in itertools.count(start=0, step=1):
        x.append(operation(i))
    return len(x)
eat_all_ram()
```

- Oh no! A misbehaving program is eating all our memory!

Interacting with Cgroups

- Create a cgroup under the memory subsystem and limit the PID to 10MB

Interacting with Cgroups

- Create a cgroup under the memory subsystem and limit the PID to 10MB

```
[root@fedora ~]# mkdir /sys/fs/cgroup/memory/dangroup
[root@fedora ~]# cd /sys/fs/cgroup/memory/dangroup
[root@fedora dangroup]# echo 10485760 > memory.limit_in_bytes
[root@fedora dangroup]# echo 2345 > tasks
```

Interacting with Cgroups

- Create a cgroup under the memory subsystem and limit the PID to 10MB

```
[root@fedora ~]# mkdir /sys/fs/cgroup/memory/dangroup
[root@fedora ~]# cd /sys/fs/cgroup/memory/dangroup
[root@fedora dangroup]# echo 10485760 > memory.limit_in_bytes
[root@fedora dangroup]# echo 2345 > tasks
```

- All fixed!

Interacting with Cgroups

- Create a cgroup under the memory subsystem and limit the PID to 10MB

```
[root@fedora ~]# mkdir /sys/fs/cgroup/memory/dangroup
[root@fedora ~]# cd /sys/fs/cgroup/memory/dangroup
[root@fedora dangroup]# echo 10485760 > memory.limit_in_bytes
[root@fedora dangroup]# echo 2345 > tasks
```

- All fixed!
- Except...process will swap near 10 MB.

Linux Namespaces

- PID
- NET
- MNT
- IPC
- UTS

Linux Namespaces

- PID
- NET
- MNT
- IPC
- UTS
- I'll only summarize PID and NET

PID Namespace

- Isolated process tree
 - eg, each tree may run its own init as pid 1
- PID Namespace is hierarchical, but one way
 - A parent namespace process is aware of child namespace processes
 - The converse is not true.
- With one parent and one child namespace, a child PID has two PID's.

NET Namespace

- Isolated network
- Separate route tables
- Separate IP space
- Separate ARP tables
- Very powerful
- Also used by OpenStack

NET Namespace

- Isolated network
- Separate route tables
- Separate IP space
- Separate ARP tables
- Very powerful
- Also used by OpenStack
- Check out ‘ip netns’

Section LXC

- LXC

LXC

- Released 2008
- First native Linux container technology
 - Prior work required kernel patches
- Powerful but limited
- Docker is to Ubuntu what LXC is to Gentoo

LXC

- LXC is lightweight, but at a cost
- Not portable between distributions
 - Docker is
- Manual configuration of networking
 - iptables, brctl, ip netns, etc
- Manual configuration of union filesystem

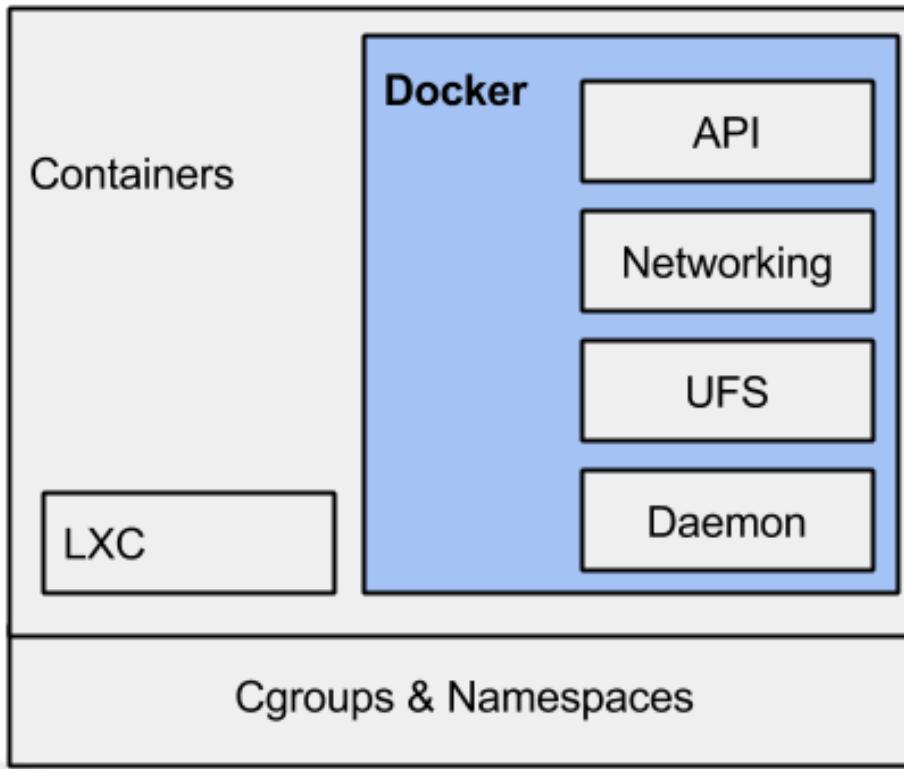
LXC versus Docker

- Both use the same cgroups, namespaces, and low level tools
- Fundamental distinction
 - Docker packages portable application centric containers
 - LXC is only a container technology

Section Docker

- Docker

Where We Are



What is Docker?

- A container implementation for Linux
- Released in 2013
- Current release 1.5
- Open Source, Apache License
 - github.com/docker/docker
- Owned by the Docker Corporation

Docker Implementation

- A monolithic daemon enabling operating system level virtualization for Linux.
- Relatively modern kernel needed
 - 2.6.32+
 - Centos/RHEL 6.5+
 - Ubuntu 12.04+
- Cgroups were merged in 2.6.24 (2008)

Platforms other than Linux?

- BSD
 - Not really
 - Jails are preferred
- Windows
 - With a Linux virtual machine
- Mac
 - With a Linux virtual machine

Platforms other than Linux?

- BSD
 - Not really
 - Jails are preferred
- Windows
 - With a Linux virtual machine
- Mac
 - With a Linux virtual machine
- Docker is a Linux container implementation

Terminology

- Image
- Container
- Layer
- Dockerfile

Image

- A set of UFS layers and metadata
- The “Ubuntu” image
- Images can be downloaded from Dockerhub
- Images are never “run”
 - Containers are
- Images are never “modified”
 - Modifications create new images

Container

- An instance of an image
- Containers are a running image with a topmost read-write UFS layer
- A container should serve one function
 - Pending business logic it can be complicated
 - Don't replicate a monolithic stack for “everything”
 - One function scales better
 - One function is easier to isolate and debug

Layer

- Docker containers are built upon UFS
- UFS is the Union File System
 - Not “strictly” a union, as upper layer can overwrite lower layer.
 - Once committed, nothing in a layer is ever deleted
- Docker images limited to 127 Layers

Dockerfile

- An image build recipe
 - Procedural
 - Basic operations for executing programs and copying files.
 - Not configuration management
- Each Dockerfile directive creates a layer
 - Maximum number of layers: 127
- Examples will follow later

Networking

- Virtual bridges
- Container Linking
 - Environment variables
 - /etc/hosts
 - Retained within Linux network namespace
- Host to container port forwarding
- iptables is used for masquerading

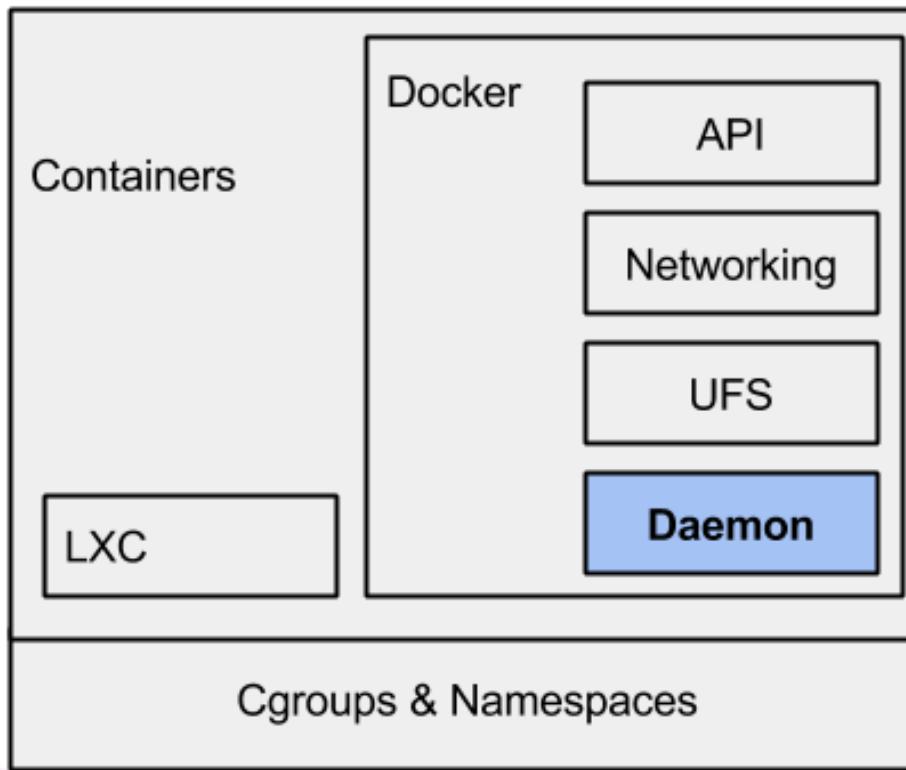
Filesystem

- Union File System
 - Base image + layers defines final image
 - Limit of 127 “layers”
 - Each docker “RUN” adds to a layer
- “Flattening” image is not natively supported
 - Achievable by exporting a container as an image
 - docker-squash utility available to help

Docker Resource Controls

- No Disk I/O Throttling
- No Disk quota
- Filesystem isolation
- CPU quotas
- RAM limits
- Network isolation but no limits

Where We Are



Docker Daemon

- Client-server Architecture
 - Monolithic design
- Daemon manages containers
- Has a REST API via HTTP and Socket
 - /var/run/docker.sock
 - HTTP will be covered later by example
- Note: API uses an Open Schema.
 - Incorrect entries are ignored.

Basic Operation

- Use the ‘docker’ command.
- This is not a guide about Docker CLI
 - Examples are for illustration only
- Download an image

```
docker pull fedora
```

- Run an image as a container

Basic Operation

- List all images

```
[root@fedora ~]# docker images
REPOSITORY          TAG        IMAGE ID      CREATED       VIRTUAL SIZE
dsklopp/multilayer  latest     fa1864ec747a  About an hour ago  229.2 MB
ubuntu              14.04     d0955f21bf24  9 days ago   188.3 MB
ubuntu              14.04.2   d0955f21bf24  9 days ago   188.3 MB
ubuntu              latest    d0955f21bf24  9 days ago   188.3 MB
ubuntu              trusty    d0955f21bf24  9 days ago   188.3 MB
ubuntu              trusty-20150320 d0955f21bf24  9 days ago   188.3 MB
<none>              <none>    00a0c78eeb6d  4 months ago  0 B
[root@fedora ~]#
```

Basic Operation

- Run an image interactively

```
docker run -t -i ubuntu:14.04 /bin/bash
```

Basic Operation

- Run an image interactively

```
docker run -t -i ubuntu:14.04 /bin/bash
```

- Start a container

```
docker start c6fda327bc33
```

Basic Operation

- Run an image interactively

```
docker run -t -i ubuntu:14.04 /bin/bash
```

- Start a container

```
docker start c6fda327bc33
```

- Stop a container

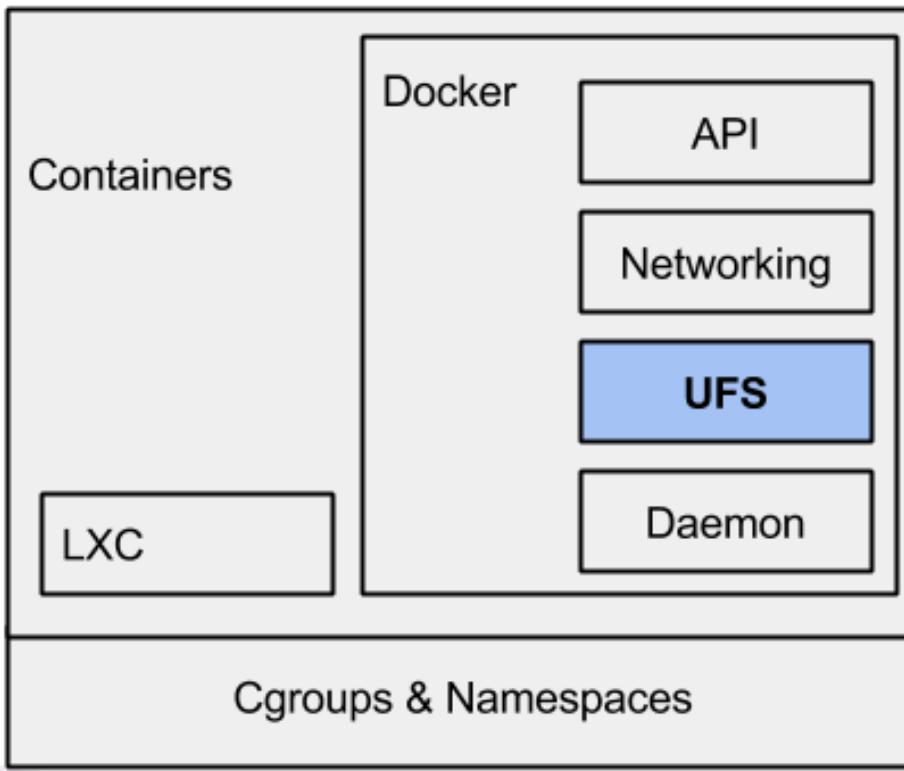
```
docker stop c6fda327bc33
```

Section

Docker Filesystem

- Docker Filesystem
- Dockerfiles
- Squashing Docker Containers

Where We Are



Union File System

- Docker utilizes Union File Systems
 - AUFS
 - BTRFS
 - VFS
 - DeviceMapper

Union File System

- Docker utilizes Union File Systems
 - AUFS
 - BTRFS
 - VFS
 - DeviceMapper
- Implementation depends on the platform
 - Fedora 21 uses DeviceMapper
 - Ubuntu 14.04 uses AUFS

Dockerfiles

- Docker UFS is best explained by Dockerfiles
- Dockerfile runs commands on an image
- Each directive creates a new layer
- The end result is an image
 - Send it to Dockerhub!
 - Share to the world!
 - Deploy it in your infrastructure

Dockerfile Structure

- Structured text file called ‘Dockerfile’

Dockerfile Structure

- Structured text file called ‘Dockerfile’
- Sample

```
FROM ubuntu:14.04
MAINTAINER Daniel Klopp

RUN touch $HOME/shell_run_variant
RUN [ "/bin/bash", "-c", "touch $HOME/exec_run_variant" ]
RUN [ "/bin/mkdir", "-p", "$HOME/exec_run_variant_mkdir" ]
RUN [ "/usr/bin/touch", "$HOME/exec_run_variant_noshell" ]
```

Dockerfile RUN

- Creates file /root/shell_run_variant

```
RUN touch $HOME/shell_run_variant
```

Dockerfile RUN

- Creates file /root/shell_run_variant

```
RUN touch $HOME/shell_run_variant
```

- Creates /root/exec_run_variant

```
RUN [ "/bin/bash", "-c", "touch $HOME/exec_run_variant" ]
```

Dockerfile RUN

- Creates file /root/shell_run_variant

```
RUN touch $HOME/shell_run_variant
```

- Creates /root/exec_run_variant

```
RUN [ "/bin/bash", "-c", "touch $HOME/exec_run_variant" ]
```

- '\$HOME/exec_run_variant_mkdir', literally.

```
RUN [ "/bin/mkdir", "-p", "$HOME/exec_run_variant_mkdir" ]
```

Dockerfile build

- Build and cropped output

```
docker build --no-cache -t dsklopp/run_variants .
```

Dockerfile build

- Build and cropped output

```
docker build --no-cache -t dsklopp/run_variants .
```

```
Step 0 : FROM ubuntu:14.04
--> d0955f21bf24
Step 1 : MAINTAINER Daniel Klopp
--> Running in 1f4b4a83bd1e
--> aa57bc7e6633
Removing intermediate container 1f4b4a83bd1e
Step 2 : RUN touch $HOME/shell_run_variant
--> Running in 24e2d896c2e5
--> 981adeaf7d60
```

The Layers of a Docker Build

- Five directives, five new layers

```
ubuntu:trusty-20150320
├─aa57bc7e6633 Virtual Size: 188.3 MB
│ └─981adeaf7d60 Virtual Size: 188.3 MB
│   └─48a88966b2b8 Virtual Size: 188.3 MB
│     └─dabec6013060 Virtual Size: 188.3 MB
│       └─8570adf87230 Virtual Size: 188.3 MB Tags: dsklopp/run_variants
```

- One layer for MAINTAINER
- Four layers for four RUN directives

Honorable Directive Mentions

- ENV, set environment variables for container
- These work in Dockerfiles, sample:

```
ENV myName="Dan Klopp" myDistro="Ubuntu"
RUN echo $myName > $HOME/name.txt
RUN echo $myDistro >> $HOME/name.txt
```

- The variables are accessible from Dockerfile and the container itself.

Honorable Directive Mentions

- Inside the container:

```
root@8ea2a330dcf1:/# env | grep -i my
myDistro=Ubuntu
myName=Dan Klopp
root@8ea2a330dcf1:/# cat $HOME/name.txt
Dan Klopp
Ubuntu
root@8ea2a330dcf1:/#
```

Honorable Directive Mentions

- Environment variables can be set on launch
- Pass to docker run command:

```
--env newKey="newValue"
```

Honorable Directive Mentions

- Environment variables can be set on launch
- Pass to docker run command:

```
--env newKey="newValue"
```

- A handy debug option!

More Directive Details

- See the official documentation
< <https://docs.docker.com/reference/builder/> >

Layers, Size and Building

- Each Layer adds to the size

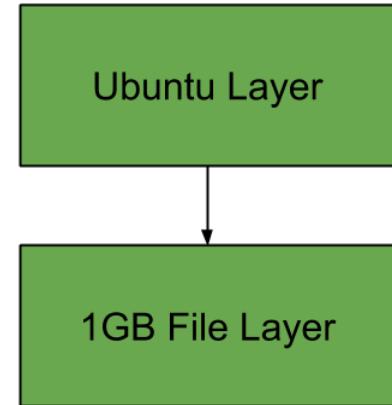
Layers, Size and Building

- Each Layer adds to the size
 - Layer 1 has Ubuntu, ~188 MB

Ubuntu Layer

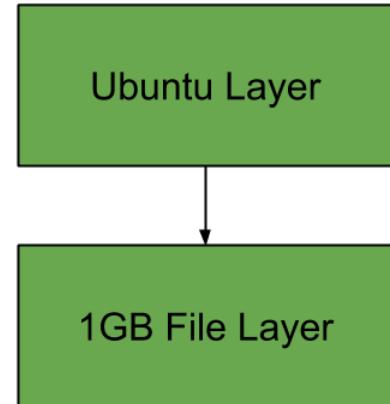
Layers, Size and Building

- Each Layer adds to the size
 - Layer 1 has Ubuntu, ~188 MB
 - Layer 2 adds a 1 GB file



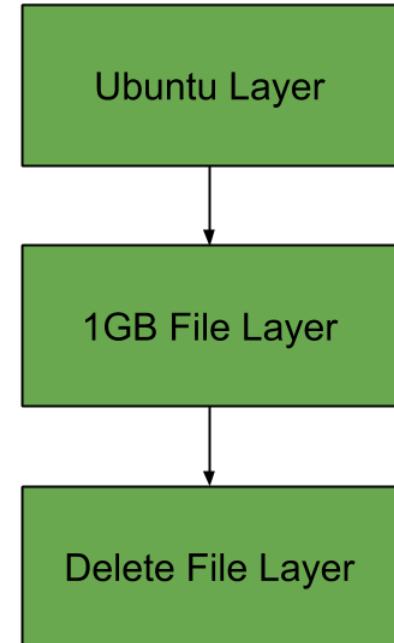
Layers, Size and Building

- Each Layer adds to the size
 - Layer 1 has Ubuntu, ~188 MB
 - Layer 2 adds a 1 GB file
 - Total size ~1.2 GB



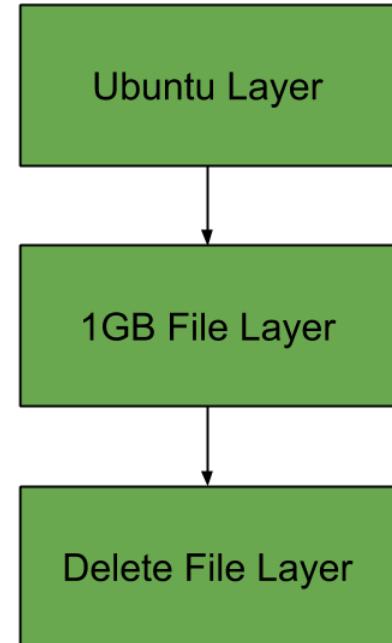
Layers, Size and Building

- Each Layer adds to the size
 - Layer 1 has Ubuntu, ~188 MB
 - Layer 2 adds a 1 GB file
 - Total size ~1.2 GB
- Delete the 1 GB file in Layer 3



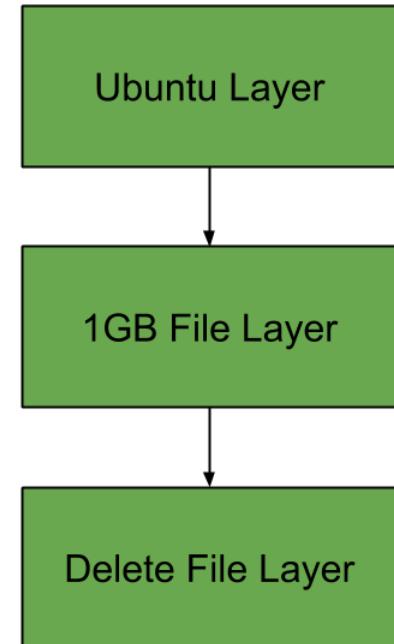
Layers, Size and Building

- Each Layer adds to the size
 - Layer 1 has Ubuntu, ~188 MB
 - Layer 2 adds a 1 GB file
 - Total size ~1.2 GB
- Delete the 1 GB file in Layer 3
 - Layer 3 deletes the 1 GB file



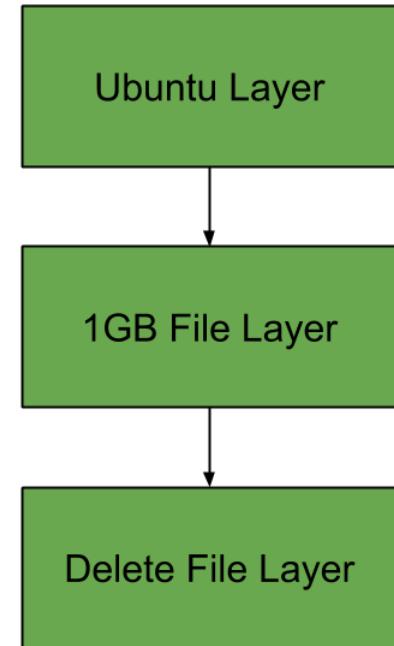
Layers, Size and Building

- Each Layer adds to the size
 - Layer 1 has Ubuntu, ~188 MB
 - Layer 2 adds a 1 GB file
 - Total size ~1.2 GB
- Delete the 1 GB file in Layer 3
 - Layer 3 deletes the 1 GB file
 - Total size ~1.2 GB



Layers, Size and Building

- Each Layer adds to the size
 - Layer 1 has Ubuntu, ~188 MB
 - Layer 2 adds a 1 GB file
 - Total size ~1.2 GB
- Delete the 1 GB file in Layer 3
 - Layer 3 deletes the 1 GB file
 - Total size ~1.2 GB
- This can be a problem



Layers, Size and Building

- Each ‘directive’ adds a layer
- Run all commands in one directive, contrast:

```
RUN dd if=/dev/zero of=/bigfile.zeros \
    bs=1k count=$((1024*1024))
RUN rm -f /bigfile.zeros
```

- With

```
RUN dd if=/dev/zero of=/bigfile.zeros \
    bs=1k count=$((1024*1024)) \
    && rm -f /bigfile.zeros
```

Layers, Size and Building

- This makes a significant difference in size

```
dsklopp/optimized    ...    188.3   MB  
dsklopp/unoptimized ...  1.262   GB
```

Dockerfiles

- Are not configuration management
- Package the results of a build process
- Each layer should be functional
 - OS layer
 - Patch to latest
 - Webserver
 - Webserver content
- A build process makes this less relevant

How To Compress Layers

- A container is a collection of layers
- In Ubuntu, this is implemented with aufs
- /var/lib/docker stores the FS and metadata.
- Edit the half dozen json files and aufs entries
- ...
- Maybe that isn't such a good idea
 - If only there was a program to do it for us
 - Jason Wilder's "docker-squash"

Docker-squash

- <https://github.com/jwilder/docker-squash>
- Trivial to use. Sample Dockerfile:

```
RUN dd if=/dev/zero of=/bigfile.zeroes \
    bs=1k count=$((1024*128))
RUN rm -f /bigfile.zeroes
RUN touch /testfile.txt
```

- Size

```
dsklopp/nonflat ... 322.5 MB
```

Docker-squash

- Present layer hierarchy

```
ubuntu:trusty, ubuntu:trusty-20150320
├─8406a4346504 Virtual Size: 188.3 MB
│ ├─76bf86be913c Virtual Size: 322.5 MB
│   ├─1f79b453eaf8 Virtual Size: 322.5 MB
│   └─d9b59ab2e0ad Virtual Size: 322.5 MB Tags: dsklopp/nonflat
```

- Compress it with docker-squash:

```
docker save d9b59ab2e0ad | \
./docker-squash --from 8406a4346504 -t dsklopp/flat | \
docker load
```

Docker-squash

- The new hierarchy:

```
ubuntu:trusty, ubuntu:trusty-20150320
├─8406a4346504 Virtual Size: 188.3 MB
│ ├─5aa152140bf7 Virtual Size: 188.3 MB Tags: dsklopp/flat:latest
│ ├─76bf86be913c Virtual Size: 322.5 MB
│   ├─1f79b453eaf8 Virtual Size: 322.5 MB
│   └─d9b59ab2e0ad Virtual Size: 322.5 MB Tags: dsklopp/nonflat
```

- Sizing:

dsklopp/flat	...	188.3	MB
dsklopp/nonflat	...	322.5	MB

Docker-squash

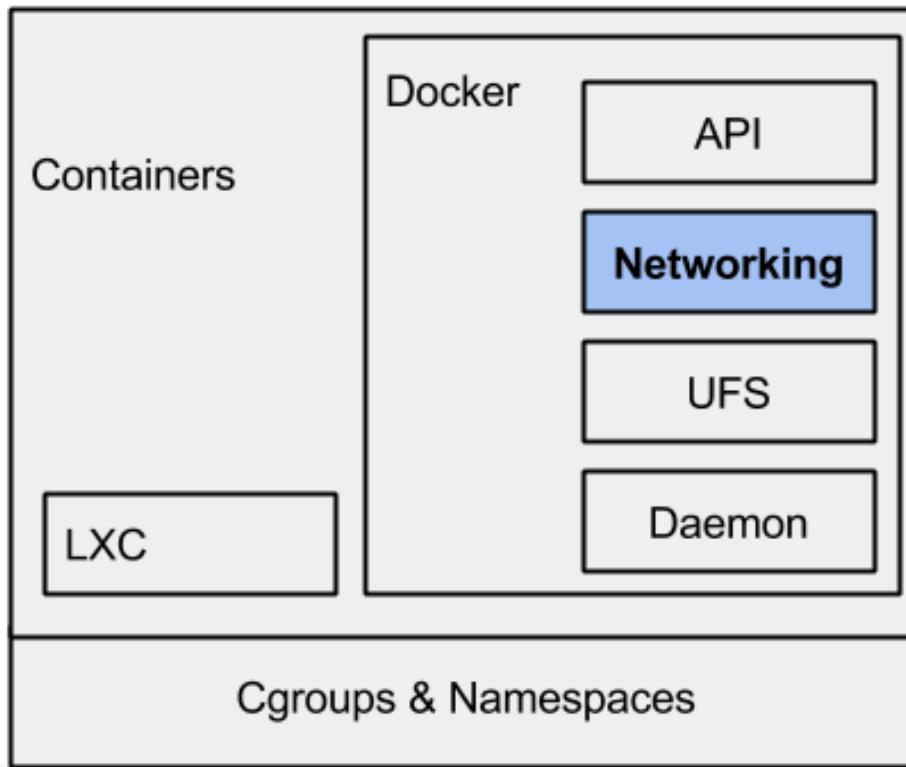
- One downside is you lose layer history
- A problem if this is significant
- If the build process is automated, it rarely is.

Section

Docker Networking

- Docker Networking

Where We Are



This Will Be Brief

- Networking deserves its own presentation to do it justice
- Docker networking is also non-trivial
- I will only go over the basics today

Default Networking

- Docker utilizes a virtual ethernet bridge
 - By default called ‘docker0’
- Containers connect via a veth interface
- Containers on docker0 can freely communicate with each other
 - Controlled by start flag ‘icc=true’, the default

Networking

- Four modes
 - None (self-explanatory)
 - Bridged (default)
 - Host
 - Shared

Bridged Networking

- Virtual Linux Ethernet bridge
- Docker default
- Dynamically or manually configured
- With `icc=true` (default), all containers can contact other containers.

Host Networking

- Docker containers use host network stack
- Though conceptually easy, not good practice
- Allows access to D-Bus
- Unexpected behavior may result
- Generally a bad idea

Shared Networking

- Existing container network stack is shared
- Each container PID and FS are distinct
- Only network namespace is shared.

Default Docker Networking

- Containers can reach out
- Outside cannot reach in without explicit port mappings
- Containers can communicate with other containers

Default Docker Networking

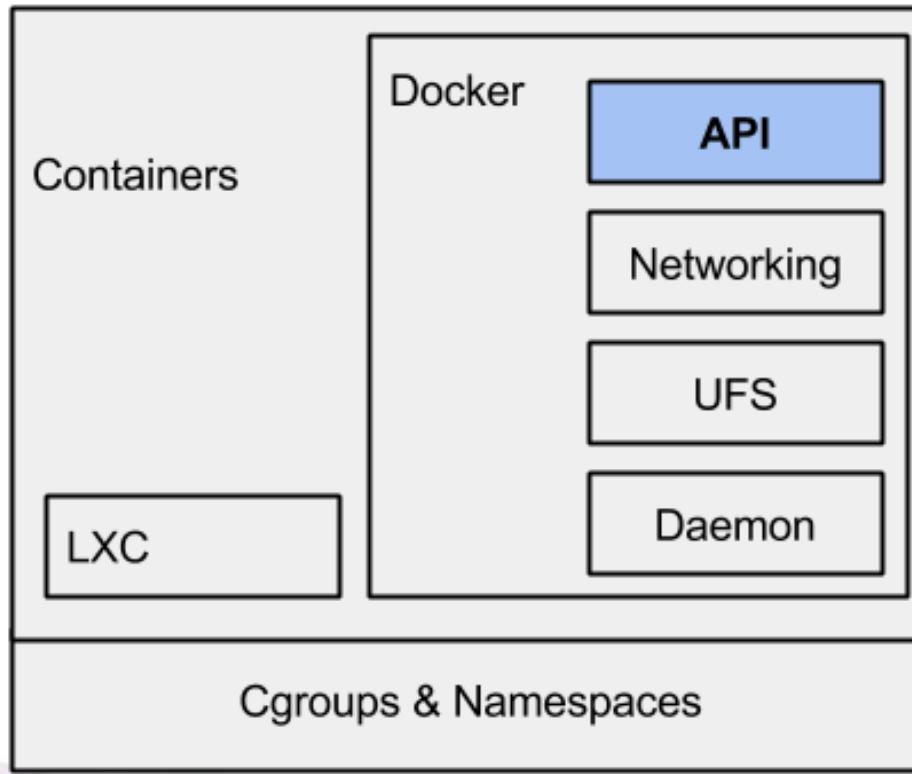
- Ultimately, it is made up of Iptables, network namespaces, and Linux bridges (virtual switches).
 - Your limits are your imagination
- If there is interest, I can create an in depth networking presentation

Section

Docker API

- Docker API

Where We Are



Docker API

- Has a (non) REST API via Unix Socket
 - /var/run/docker.sock
- Note: API uses an Open Schema.
 - Incorrect entries are ignored.
- There is an HTTP endpoint
 - disabled by default

Docker API via HTTP

- Disabled by default
- Pass “-H \$host:\$port” to daemon
 - Ubuntu: /etc/default/docker
 - Fedora: /etc/sysconfig/docker

Docker API via HTTP

- Disabled by default
- Pass “-H \$host:\$port” to daemon
 - Ubuntu: /etc/default/docker
 - Fedora: /etc/sysconfig/docker
- It is REST-like, but not REST
 - It does not allow querying for URI's
 - Certain operations are hijacked for STDOUT/ERR

Docker API via HTTP

- Disabled by default
- Pass “-H \$host:\$port” to daemon
 - Ubuntu: /etc/default/docker
 - Fedora: /etc/sysconfig/docker
- It is REST-like, but not REST
 - It does not allow querying for URI's
 - Certain operations are hijacked for STDOUT/ERR
- https://docs.docker.com/reference/api/docker_remote_api

Docker API Query

```
curl -sX GET 127.0.0.1:2375/info | python -m json.tool
```

Docker API Query

```
curl -sX GET 127.0.0.1:2375/info | python -m json.tool
```

```
{  
    "Containers": 5,  
    "Debug": 0,  
    "DockerRootDir": "/var/lib/docker",  
    "Driver": "devicemapper",  
    "DriverStatus": [  
        [  
            "Pool Name",  
            "docker-253:1-263795-pool"  
        ],  
        [  
            "Status",  
            "Docker daemon is running"  
        ]  
    ],  
    "Events": 0,  
    "GraphDriver": "devicemapper",  
    "GraphOptions": {},  
    "HostConfig": {},  
    "Info": {  
        "APIVersion": "1.35",  
        "Author": "Docker, Inc.",  
        "Build": "1705.03-df62cde",  
        "Name": "Docker",  
        "OperatingSystem": "Debian 9 (Stretch)",  
        "ServerVersion": "17.05.0-ce",  
        "Version": "17.05.0-ce"  
    },  
    "MountLabel": "",  
    "NFSLabels": {},  
    "NetworkSettings": {},  
    "PluginStatus": {},  
    "Plugins": {},  
    "Status": "running",  
    "TLSInfo": {}  
}
```

Docker API

- Queries are nice, but wouldn't it be awesome if we could create containers via the API?

Docker API

- Queries are nice, but wouldn't it be awesome if we could create containers via the API?
 - We can

Docker API

- Queries are nice, but wouldn't it be awesome if we could create containers via the API?
 - We can
 - POST /containers/create

Docker API

- POST /containers/create
- That's easy!

Docker API

- POST /containers/create
- That's easy!
- Pet peeve: online documentation for generating proper POST requests is scant

Docker API Launch Container

- The Docker POST requires a configuration.
- Get one from a running container
- Start an interactive container

```
docker run -t -i ubuntu:14.04 /bin/bash
```

Docker API Launch Container

- Get the container's ID

```
curl -sX GET 127.0.0.1:2375/containers/json
```

Docker API Launch Container

- Get the container's ID

```
curl -sX GET 127.0.0.1:2375/containers/json
```

```
[  
  {  
    "Command": "/bin/bash",  
    "Created": 1427597350,  
    "Id": "027912d4100bebc3a43bdd1c466df2142ac0cda2bddc876",  
    "Image": "ubuntu:14.04",  
    "Names": [  
      "/angry_pike"  
    ],  
    "Ports": [],  
    "Status": "Up 32 seconds"  
  }]  
]
```

Docker API Launch Container

- Get the container's Configuration

```
curl -X GET 127.0.0.1:2375/containers/$ID/json
```

- Take the config output (cropped below)

```
"Config": {  
    "AttachStderr": true,  
    "AttachStdin": true,  
    "AttachStdout": true,  
    "Cmd": [  
        "/bin/bash"  
    ],
```

- Docker inspect will also display this information.

Docker API Launch Container

- Modify the config to your liking, and send it as the POST body. For example

```
curl -sX POST -H "Content-Type: application/json" \
http://127.0.0.1:2375/containers/create -d '{ $CONTENT_HERE }'
```

- Where \$CONTENT_HERE is the POST body you copied from before.
- Cropped output:

```
{"Id": "ae821e145e6f1e9beb4848722d2fde82d33535756c2a9625430bb1d88}
```

Docker API Launch Container

- The container isn't started yet.
- Start it with the ID returned from the POST

```
curl -sX POST -H 'Content-Type: application/json'  
127.0.0.1:2375/containers/$IDNEW/start
```

- Check on your container

```
docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
ae821e145e6f	ubuntu:14.04	"/bin/bash"	4 minutes ago	Up 2 sec

Docker API

- Start / stop containers
- Import / export images
- Look at running processes
- Attach to a running container
- Extremely powerful
 - Be wary of access, it is inherently insecure

Section

Docker Management

- Docker Management

Docker Management

- In short, there aren't.

Docker Management

- In short, there aren't.
- Ok, that was harsh.
 - This is a very new field.
 - There are a lot of newfound projects for this gap

Docker Management

- A major gap
- New field, largely untested, mostly custom
- Noteworthy Solutions
 - Swarm (official)
 - Machine
 - Compose
 - Mesos

Docker Management

- Config Management is entering the race too!
 - Chef
 - Puppet
 - Ansible

Summary

- Docker is built upon Linux Containers
 - Cgroups
 - Namespaces
 - Union File Systems
- Docker provides single-host management functionality of containers
- Containers on one system will work on another.

What We Covered

- Why containers are important
- Kernel features enabling containerization
 - cgroups
 - namespaces
- LXC and its complexities

What We Covered

- Docker Implementation
 - General operation
 - Filesystem
 - Networking
 - API
- Docker Management, or lack thereof

Summary

- We talked about the fishing rod
 - Now go fish!

Special Thanks

- Taos
- SVLUG
- GE
- Especially
 - Rachel Pecchenino
 - Kevin Dankwardt

Questions?